

ST509 Computational Statistics

Lecture 12: Gibbs Samplers

Seung Jun Shin

Department of Statistics
Korea University

E-mail: `sjshin@korea.ac.kr`



Gibbs Sampler I

- ▶ We extend the idea of MCMC by breaking the large complex problems into a sequence of simpler ones.
- ▶ **Gibbs sampler** may take a long time to converge, but is an interesting candidate.

Gibbs Sampler II

- Suppose we like to sample $(X, Y) \sim f(x, y)$, the two-stage Gibbs sampler generates a Markov chain (X_t, Y_t) whose stationary distribution is $f(x, y)$ as follows.

1. Take $X_0 = x_0$
2. For $t = 1, 2, \dots$, generate

$$2.1 \quad Y_t \sim f_{Y|X}(\cdot \mid x_{t-1})$$

$$2.2 \quad X_t \sim f_{X|Y}(\cdot \mid y_t)$$

- Straightforward to implement as long as simulating from both conditionals are feasible.

Gibbs Sampler III

ex1 Consider

$$(X, Y) \sim N_2 \left(0, \begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix} \right)$$

- ▶ Given x_t , generate

$$Y_{t+1} \mid x_t \sim N(\rho x_t, 1 - \rho^2)$$

$$X_{t+1} \mid y_{t+1} \sim N(\rho y_{t+1}, 1 - \rho^2)$$

- ▶ Please implement the Gibbs sampler to generate the samples.

Gibbs Sampler IV

ex2 Consider the pair of distributions

$$X \mid \theta \sim \text{Binomial}(n, \theta), \quad \theta \sim \text{Beta}(a, b)$$

which leads to the joint distribution:

$$f(x, \theta) = \binom{n}{x} \frac{\Gamma(a+b)}{\Gamma(a)\Gamma(b)} \theta^{x+a-1} (1-\theta)^{n-x+b-1}$$

- Notice that the distribution of $X \mid \theta$ is given and

$$\theta \mid x \sim \text{Beta}(x+a, n-x+b)$$

- Please implement the Gibbs sampler to generate the samples.

Gibbs Sampler V

- ▶ Extension to multi-stage Gibbs sampler is straightforward.
 1. At iteration $t = 1, 2, \dots$, given $\mathbf{x}^{(t)} = (x_1^{(t)}, \dots, x_p^{(t)})$ generate
 - 1.1 $X_1^{(t+1)} \sim f_1(x_1 \mid x_2^{(t)}, \dots, x_p^{(t)})$
 - 1.2 $X_2^{(t+1)} \sim f_2(x_2 \mid x_1^{(t+1)}, x_3^{(t)}, \dots, x_p^{(t)})$
 - \vdots
 - 1.3 $X_p^{(t+1)} \sim f_p(x_p \mid y_t)x_1^{(t+1)}, \dots, x_{p-1}^{(t+1)})$
- ▶ We call f_1, \dots, f_p the **full conditionals**, the only densities used for simulation.

Gibbs Sampler VI

- Bayesian Linear Model:

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i, \quad N(0, \sigma^2)$$

- We consider a conjugate priors:

$$\beta_j \sim N(\mu_j, \tau_j^2), \quad j = 0, 1 \quad \text{and} \quad \sigma^2 \sim 1/\text{Gamma}(a, b)$$

- Gibbs sampler can be used to generate sample from

$$f(\beta_0, \beta_1, \sigma^2 \mid \mathbf{y}) \propto \pi_1(\beta_0) \pi_1(\beta_1) \pi_2(\sigma^2) f(\mathbf{y} \mid \beta_0, \beta_1, \sigma^2)$$

Gibbs Sampler VII

- It is not difficult to show that the full conditionals are given by

$$\begin{aligned}\beta_0 \mid \beta_1, \sigma^2, \mathbf{y} &\sim N \left(\frac{\sum_{i=1}^n (y_i - \beta_1 x_i)^2 / \sigma^2 + \mu_0 / \tau_0^2}{n / \sigma^2 + 1 / \tau_0^2}, (n / \sigma^2 + 1 / \tau_0^2)^{-1} \right) \\ \beta_1 \mid \beta_0, \sigma^2, \mathbf{y} &\sim N \left(\frac{\sum_{i=1}^n x_i (y_i - \beta_0) / \sigma^2 + \mu_1 / \tau_1^2}{\sum_{i=1}^n x_i^2 / \sigma^2 + 1 / \tau_1^2}, \left(\sum_{i=1}^n x_i^2 / \sigma^2 + 1 / \tau_1^2 \right)^{-1} \right) \\ \sigma^2 \mid \beta_0, \beta_1, \mathbf{y} &\sim 1 / \text{Gamma} \left(\frac{n}{2} + a, \frac{1}{2} \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2 + b \right)\end{aligned}$$

Gibbs Sampler VIII

```
# Gibbs_sampler
for(i in 1:n.samples){
  # update sigma2:
  SSE <- sum((y - beta[1] - x * beta[2])^2)
  sigma2 <- 1/rgamma(1, n/2 + a, SSE/2 + b)

  # update beta1:
  v <- n/sigma2 + 1/tau[1]^2
  m <- sum(y - x * beta[2])/sigma2 + mu[1]/tau[1]^2
  beta[1] <- rnorm(1, m/v, 1/sqrt(v))

  # update beta2:
  v <- sum(x^2)/sigma2 + 1/tau[2]^2
  m <- sum(x*(y-beta[1]))/sigma2 + mu[2]/tau[2]^2
  beta[2] <- rnorm(1,m/v,1/sqrt(v))

  samples[i,] <- c(beta, sigma2)
}
```

MH Algorithm within Gibbs Sampler I

- ▶ Full conditionals are not always available
- ▶ One can use MH algorithm to get a sample from a targeted conditional density.

MH Algorithm within Gibbs Sampler II

- Bayesian Logistic Regression

$$y_i \mid \mathbf{x}_i \sim \text{Bern}(p(\mathbf{x}_i; \boldsymbol{\beta}_0))$$

where $\mathbf{x}_i = (1, x_{i1}, \dots, x_{ip})^T$ and $\boldsymbol{\beta} = (\beta_0, \beta_1, \dots, \beta_p)^T$.

- Independent normal priors on $\boldsymbol{\beta} = (\beta_0, \beta_1, \dots, \beta_p)^T$:

$$\beta_j \sim N(\mu_j, \tau_j^2), \quad j = 0, 1, \dots, p$$

- The full conditional density of β_j given all others is

$$f(\beta_j \mid \boldsymbol{\beta}_{-j}, \mathbf{y}) \propto \underbrace{\prod_{i=1}^n \{p(\mathbf{x}_i; \boldsymbol{\beta})\}_i^y \{1 - p(\mathbf{x}_i; \boldsymbol{\beta})\}^{1-y_i}}_{\text{Likelihood}} \times \underbrace{\exp\left(-\frac{1}{2\tau_j^2}(\beta_j - \mu_j)^2\right)}_{\text{Prior}}.$$

- Apply MH to get samples from the full conditional.

MH Algorithm within Gibbs Sampler III

```
mcmc.logit2 <- function(x, y, init, n.sample = 10000, step = rep(0.3, p)){  
  n <- nrow(x)  
  p <- ncol(x)  
  
  post.beta <- matrix(0, n.sample, p)  
  ac.ratio <- matrix(0, n.sample, p)  
  
  prior.m <- 10  
  prior.s <- 1000 # for vague prior  
  
  # intialize  
  post.beta[1,] <- beta <- init  
  eta <- x %*% beta  
  pi <- exp(eta)/(1 + exp(eta))  
  
  log.like <- sum(y * log(pi) + (1 - y) * log(1 - pi))  
  
  for (iter in 1:n.sample){  
    beta.new <- beta  
    # candidate  
    for (j in 1:p)  
    {  
      beta.new[j] <- beta[j] + rnorm(1, 0, step[j])  
      eta.new <- x %*% beta.new  
      pi.new <- exp(eta.new)/(1 + exp(eta.new))
```

MH Algorithm within Gibbs Sampler IV

```
# prior
log.prior      <- dnorm(beta[j],      prior.m, prior.s, log = T)
log.prior.new  <- dnorm(beta.new[j], prior.m, prior.s, log = T)

# likelihood
log.like.new  <- sum(y * log(pi.new) + (1 - y) * log(1 - pi.new))

# ratio
temp <- exp((log.like.new + log.prior.new) - (log.like + log.prior))
rho <- min(1, temp)

if (runif(1) < rho) {
  ac.ratio[iter,j] <- 1
  beta[j] <- beta.new[j]
  log.like <- log.like.new
  eta <- x %*% beta
  pi <- exp(eta)/(1 + exp(eta))
}

post.beta[iter,] <- beta
}

obj <- list(posterior = post.beta, acpt.ratio = apply(ac.ratio, 2, mean))
return(obj)
}
```

Nonparametric Function Estimation I

- ▶ Consider

$$y_i = f(x_i) + \epsilon_i$$

with a random error ϵ_i and $E(\epsilon_i) = 0$ and $\text{Var}(\epsilon_i) = \sigma^2$.

- ▶ Suppose $x \in [0, 1]$. Bernstein polynomials of f is given by

$$\begin{aligned} B_M(x; f) &= \sum_{m=0}^M f(m/M) \binom{M}{m} x^m (1-x)^{M-m} \\ &= \sum_{m=0}^M \beta_m b_M(x, m) \end{aligned}$$

- ▶ Weierstrass approximation theorem states that

$$\lim_{M \rightarrow \infty} \sup_x |f(x) - B_M(x; f)| = 0.$$

- ▶ That is, we can model $f(x)$ via Bernstein polynomials.

Nonparametric Function Estimation II

- Consider a re-parameterization of $\beta_m = f(m/M)$, $m = 0, 1, \dots, M$.

$$\gamma_0 = \beta_0, \gamma_m = \beta_m - \beta_{m-1}, \quad j = 1, \dots, M-1$$

which yields

$$\beta_m = \sum_{k=0}^m \gamma_k.$$

- Notice that

$$\begin{aligned} B_M(x; f) &= \sum_{m=0}^M \beta_m b_M(x, m) \\ &= \sum_{m=0}^M \sum_{k=0}^m \gamma_k b_M(x, m) \\ &= \sum_{k=0}^M \gamma_k \sum_{m=k}^M b_M(x, m) = \sum_{k=0}^M \gamma_k F_M(x, k) \end{aligned}$$

Nonparametric Function Estimation III

where

$$F_M(x, k) = \sum_{m=k}^M b_M(x, k) = \int_0^x \frac{u^{m-1}(1-u)^{M-k}}{\text{Beta}(k, M-k+1)} du$$

i.e., the distribution function of the Beta random variable with parameters m and $M - m + 1$ for $m = 0, 1, 2, \dots, M$.

- ▶ Bernstein polynomial model for $f(x)$ is given by

$$y_i \sim N \left(\sum_{k=0}^M \gamma_k F_M(x_i, k), \sigma^2 \right)$$

- ▶ One can assume the priors on $\gamma_j, j = 0, \dots, M$ and σ^2 .

Nonparametric Function Estimation IV

- ▶ Bernstein polynomial is particularly useful under presence of shape constraints.
- ▶ We consider a nonparametric dose response model

$$y_i \mid x_i \sim \text{Bern}(p(x_i))$$

where

$$p(x) = P(Y = 1 \mid x) \in [0, 1]$$

- ▶ Here $p(x)$ is assumed to be an increasing function of x and $p(0) = 0$ and $p(1) \leq 1$, with $x \in [0, 1]$.

Nonparametric Function Estimation V

► The assumptions can be imposed by:

1. $\gamma_0 = 0$,
2. $\gamma_j \geq 0, \quad j = 1, \dots, p$;
3. $\beta_M = \sum_{m=1}^M \gamma_m \leq 1$

Nonparametric Function Estimation VI

- ▶ Applying the Bernstein approximation, we have

$$y_i \mid x_i \sim \text{Bern}(p(x))$$

where

$$p(x) = \sum_{k=1}^M \gamma_k F_M(x_i, k)$$

- ▶ To impose the 2nd and 3rd constraints, we consider

$$\gamma_1 = \delta_1, \quad \text{and} \quad \gamma_j = \delta_j \prod_{k=1}^{j-1} (1 - \delta_k)$$

with the independent uniform prior:

$$\delta_j \stackrel{iid}{\sim} \text{Uniform}(0, 1), \quad j = 1, 2, \dots, p.$$

Nonparametric Function Estimation VII

```
set.seed(1)
n <- 200
x <- runif(n)

beta <- 5
f <- function(x, beta) (exp(beta * x) - 1)/(exp(beta) - 1)
p <- f(x, beta)
y <- rbinom(n, 1, p)

n.sample <- 10000
step <- 0.1
M <- 10 # degrees of bernstein polynomials

max.x <- 1
min.x <- 0
x.tilde <- (x - min.x)/(max.x - min.x)
W <- unlist(lapply(1:M, function(k) pbeta(x.tilde, k, M-k+1)))
Ft <- matrix(W, ncol = M)

# initialize
delta <- runif(M)

# compute gamma from delta
gamma <- delta * cumprod(1 - c(0, delta[-M]))

# p from gamma
p <- Ft %*% gamma
```

Nonparametric Function Estimation VIII

```
# likelihood
log.like <- sum(y * log(p) + (1-y) * log(1-p))

ac.ratio <- matrix(0, n.sample, M)
post.gamma <- matrix(0, n.sample, M)

for (iter in 1:n.sample) {
  delta.new <- delta
  gamma.new <- gamma
  for (m in 1:M)
  {
    eta <- log(delta[m]/(1 - delta[m]))
    eta.new <- eta + rnorm(1, 0, step)
    delta.new[m] <- exp(eta.new)/(1 + exp(eta.new))

    gamma.new <- delta.new * cumprod(1 - c(0, delta.new[-M]))

    # update p
    p.new <- Ft %*% gamma.new

    # likelihood
    log.like.new <- sum(y * log(p.new) + (1-y) * log(1-p.new))

    # update
    temp <- exp(log.like.new - log.like)
    rho <- min(1, temp)
    if (runif(1) < rho) {
      ac.ratio[iter,m] <- 1
    }
  }
}
```

Nonparametric Function Estimation IX


```
      delta[m] <- delta.new[m]
      gamma[m] <- gamma.new[m]
      log.like <- log.like.new
      p <- Ft %*% gamma
    }
  }
  post.gamma[iter,] <- gamma
}

est <- apply(Ft %*% t(post.gamma), 1, quantile, prob = c(0.025, 0.5, 0.975))


id <- order(x)

plot(x[id], est[2, id],
     xlab = "x", ylab = "probability", col = 4, type = "l", lwd = 2,
     ylim = c(0,1))
polygon(c(x[id], rev(x[id])), c(est[1, id], rev(est[3, id])),
       col = rgb(0,0,1, alpha = 0.22),
       border = F)
t <- seq(0, 1, length = 1000)
lines(t, f(t, beta), col = 2, lwd = 2, lty = 2)
```

Nonparametric Function Estimation X



(a) $\beta = 1$



(b) $\beta = 5$

Figure: Estimated Curves with 95% credible bands. Dashed line is the true and solid line is the estimated.