

ST720 Data Science

Data Transformation with dplyr

Seung Jun Shin (sjshin@krea.ac.kr)

Department of Statistics, Korea University

package : tidyverse

```
library(tidyverse)
library(nycflights13)
```

- ▶ Illustrate the dplyr package to flights data.
- ▶ **dplyr** overwrites some functions in base R.
- ▶ Use their full names: `stats::filter()` and `stats::lag()`.

flights data in nycflights13

```
flights
```

```
## # A tibble: 336,776 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517           515           2     830
## 2  2013     1     1     533           529           4     850
## 3  2013     1     1     542           540           2     923
## 4  2013     1     1     544           545          -1    1004
## 5  2013     1     1     554           600          -6     812
## 6  2013     1     1     554           558          -4     740
## 7  2013     1     1     555           600          -5     913
## 8  2013     1     1     557           600          -3     709
## 9  2013     1     1     557           600          -3     838
## 10 2013     1     1     558           600          -2     753
```

```
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <
## #   minute <dbl>, time_hour <dtm>
```

nycflights13

- ▶ The data looks a little different from the data frame.
- ▶ To see the whole data set, `view(flights)`.
- ▶ Abbreviations under the column names describe the variable type.
 - ▶ `int` for integer.
 - ▶ `dbl` for doubles or real numbers.
 - ▶ `chr` for characters or strings.
 - ▶ `dtm` for date-times.
 - ▶ `lgl` for logical values.
 - ▶ `fctr` for factors.
 - ▶ `date` for dates.

dplyr Basics

- ▶ `filter()`: Pick observations by their value.
- ▶ `select()`: Pick variables by their names.
- ▶ `mutate()`: Create new variables with functions of existing variables.
- ▶ `arrange()`: Reorder the rows.
- ▶ `summarize()`: Collapse many values down to a single summary.
- ▶ All functions can be used in conjunction with `group_by()`.

dplyr Basics

- ▶ First argument is a data frame.
- ▶ Next ones describes what to do with dhata frame, using variable names (without quotes).
- ▶ Output is a new data frame.

Filter Rows with filter()

```
jan1 <- filter(flights, month == 1, day == 1)
jan1
```

```
## # A tibble: 842 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013     1     1     517           515         2     830
## 2  2013     1     1     533           529         4     850
## 3  2013     1     1     542           540         2     923
## 4  2013     1     1     544           545        -1    1004
## 5  2013     1     1     554           600        -6     812
## 6  2013     1     1     554           558        -4     740
## 7  2013     1     1     555           600        -5     913
## 8  2013     1     1     557           600        -3     709
## 9  2013     1     1     557           600        -3     838
## 10 2013     1     1     558           600        -2     753
```

```
## # ... with 832 more rows, and 12 more variables: sched_arr_time <int>
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <
## #   minute <dbl>, time_hour <dtm>
```

Logical Operations

- ▶ Equivalent code 1

```
filter(flights, month == 11 | month == 12)  
filter(flights, month %in% c(11,12))
```

- ▶ Equivalent code 2

```
filter(flights, !(arr_delay > 120 | dep_delay > 12))  
filter(flights, arr_delay <= 120, dep_delay <= 12)
```

- ▶ `is.na()` and `near()` functions are useful.

Select Columns with `select()`

```
select(flights, year, month, day)
select(flights, year:day)
select(flights, -(year:day))
```

There are some useful functions you can use within `select()`.

- ▶ `starts_with("abc")`: matches names that begin with "abc".
- ▶ `ends_with("xyz")`: matches names that end with "xyz".
- ▶ `contains("ijk")`: matches names that contain "ijk".
- ▶ `num_range("x", 1:3)`: matches x1, x2, and x3

Select Columns with select()

- ▶ rename() is useful to change the variable name while keeping all other variables.

```
rename(flights, tail_num = tailnum)
```

```
## # A tibble: 336,776 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517           515           2     830
## 2  2013     1     1     533           529           4     850
## 3  2013     1     1     542           540           2     923
## 4  2013     1     1     544           545          -1    1004
## 5  2013     1     1     554           600          -6     812
## 6  2013     1     1     554           558          -4     740
## 7  2013     1     1     555           600          -5     913
## 8  2013     1     1     557           600          -3     709
## 9  2013     1     1     557           600          -3     838
## 10 2013     1     1     558           600          -2     753
```

```
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tail_num <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <
## #   minute <dbl>, time_hour <dtm>
```

Select Columns with select()

- To move some variables to the front, use everything() helper.

```
select(flights, time_hour, air_time, everything())
```

```
## # A tibble: 336,776 x 19
```

```
##   time_hour          air_time  year month   day dep_time sched_dep
##   <dtm>              <dbl> <int> <int> <int>   <int>
## 1 2013-01-01 05:00:00      227  2013     1     1     517
## 2 2013-01-01 05:00:00      227  2013     1     1     533
## 3 2013-01-01 05:00:00      160  2013     1     1     542
## 4 2013-01-01 05:00:00      183  2013     1     1     544
## 5 2013-01-01 06:00:00      116  2013     1     1     554
## 6 2013-01-01 05:00:00      150  2013     1     1     554
## 7 2013-01-01 06:00:00      158  2013     1     1     555
## 8 2013-01-01 06:00:00       53  2013     1     1     557
## 9 2013-01-01 06:00:00      140  2013     1     1     557
## 10 2013-01-01 06:00:00      138  2013     1     1     558
## # ... with 336,766 more rows, and 12 more variables: dep_delay <dbl>
## #   arr_time <int>, sched_arr_time <int>, arr_delay <dbl>, carrier <
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>, distance
## #   hour <dbl>, minute <dbl>
```

Add new variables with `mutate()`

- ▶ Useful to add new columns that are functions of existing columns.
- ▶ `mutate()` always adds new columns at the end of the data set.
- ▶ Create a small data set with less variables:

```
flights_sml <- select(flights, year:day,  
                      ends_with("delay"),  
                      distance, air_time)
```

Add new variables with mutate()

```
mutate(flights_sml,  
       gain = arr_delay - dep_delay,  
       speed = distance / air_time)
```

```
## # A tibble: 336,776 x 9
```

```
##   year month   day dep_delay arr_delay distance air_time  gain speed  
##   <int> <int> <int>     <dbl>     <dbl>     <dbl>     <dbl> <dbl> <dbl>  
## 1  2013     1     1         2         11     1400      227     9  6.1  
## 2  2013     1     1         4         20     1416      227    16  6.1  
## 3  2013     1     1         2         33     1089      160    31  6.1  
## 4  2013     1     1        -1        -18     1576      183   -17  8.1  
## 5  2013     1     1        -6        -25      762      116   -19  6.1  
## 6  2013     1     1        -4         12      719      150    16  4.1  
## 7  2013     1     1        -5         19     1065      158    24  6.1  
## 8  2013     1     1        -3        -14      229       53   -11  4.1  
## 9  2013     1     1        -3         -8      944      140    -5  6.1  
## 10 2013     1     1        -2          8      733      138    10  5.1  
## # ... with 336,766 more rows
```

Add new variables with mutate()

- You can refer to columns that you've just created.

```
mutate(flights_sml,  
  gain = arr_delay - dep_delay,  
  hours = air_time/60,  
  gain_per_hour = gain/hours)
```

```
## # A tibble: 336,776 x 10
```

```
##   year month   day dep_delay arr_delay distance air_time  gain hours  
##   <int> <int> <int>     <dbl>     <dbl>     <dbl>     <dbl> <dbl> <dbl>  
## 1  2013     1     1         2         11      1400      227     9 3.7  
## 2  2013     1     1         4         20      1416      227    16 3.7  
## 3  2013     1     1         2         33      1089      160    31 2.6  
## 4  2013     1     1        -1        -18      1576      183   -17 3.0  
## 5  2013     1     1        -6        -25       762      116   -19 1.9  
## 6  2013     1     1        -4         12       719      150    16 2.5  
## 7  2013     1     1        -5         19      1065      158    24 2.6  
## 8  2013     1     1        -3        -14       229       53   -11 0.8  
## 9  2013     1     1        -3         -8       944      140    -5 2.3  
## 10 2013     1     1        -2          8       733      138    10 2.3  
## # ... with 336,766 more rows, and 1 more variable: gain_per_hour <dbl>
```

Add new variables with `transmute()`

- If you only want to keep the new variables, use `transmute()`.

```
transmute(flights_sml,  
  gain = arr_delay - dep_delay,  
  hours = air_time/60,  
  gain_per_hour = gain/hours)
```

```
## # A tibble: 336,776 x 3  
##   gain hours gain_per_hour  
##   <dbl> <dbl>         <dbl>  
## 1      9 3.78           2.38  
## 2     16 3.78           4.23  
## 3     31 2.67          11.6  
## 4    -17 3.05          -5.57  
## 5    -19 1.93          -9.83  
## 6     16 2.5            6.4  
## 7     24 2.63           9.11  
## 8    -11 0.883         -12.5  
## 9     -5 2.33          -2.14  
## 10    10 2.3            4.35  
## # ... with 336,766 more rows
```

Useful Creation Functions

- ▶ There are many functions you can use with `mutate()`.
- ▶ Key is that the function must be vectorized: Both input and output should be vectors.
 - ▶ Arithmetic Operators
 - ▶ Modular Arithmetic
 - ▶ Logs
 - ▶ Offsets such as `lead()`, `lag()`
 - ▶ Cumulative and rolling aggregates
 - ▶ Logical Comparisons
 - ▶ Ranking

arrange()

- ▶ `sort()` will sort a vector, but not a data frame.
- ▶ `arrange()` is for it.
- ▶ Specify the data frame and the column by which you want it to be sorted.

```
arrange(flights_sml, distance)
```

```
## # A tibble: 336,776 x 7
##   year month   day dep_delay arr_delay distance air_time
##   <int> <int> <int>     <dbl>     <dbl>     <dbl>     <dbl>
## 1  2013     7    27         NA         NA         17         NA
## 2  2013     1     3         -2         -2         80         30
## 3  2013     1     4         40         27         80         30
## 4  2013     1     4        134        136         80         28
## 5  2013     1     4         -1         -6         80         32
## 6  2013     1     5         -5        -25         80         29
## 7  2013     1     6         -4          0         80         22
## 8  2013     1     7         -5        -12         80         25
## 9  2013     1     8         -3         39         80         30
## 10 2013     1     9         -3         -7         80         27
## # ... with 336,766 more rows
```

arrange()

- To break ties, we can further sort by additional variables.

```
arrange(flights_sml, distance, year, month, day)
```

```
## # A tibble: 336,776 x 7
```

```
##   year month   day dep_delay arr_delay distance air_time
##   <int> <int> <int>    <dbl>    <dbl>    <dbl>    <dbl>
## 1  2013     7    27         NA         NA        17         NA
## 2  2013     1     3         -2         -2        80         30
## 3  2013     1     4         40         27        80         30
## 4  2013     1     4        134        136        80         28
## 5  2013     1     4         -1         -6        80         32
## 6  2013     1     5         -5        -25        80         29
## 7  2013     1     6         -4          0        80         22
## 8  2013     1     7         -5        -12        80         25
## 9  2013     1     8         -3         39        80         30
## 10 2013     1     9         -3         -7        80         27
## # ... with 336,766 more rows
```

Grouped Summaries with `summarize()`

- ▶ `summarize()` collapses a data frame to a single row.

```
summarize(flights, N = n(), delay = mean(dep_delay, na.rm = TRUE))
```

```
## # A tibble: 1 x 2
##       N delay
##   <int> <dbl>
## 1 336776 12.6
```

- ▶ You can specify a list of functions to summarize the observations.

Grouped Summaries with `summarize()`

- ▶ `summarize()` is much useful with `group_by()`

```
by_day <- group_by(flights, year, month, day)
```

- ▶ `by_day` is now grouped data frame.
- ▶ Following code yields grouped summaries.

```
summarize(by_day, N = n(), delay = mean(dep_delay, na.rm = FALSE))
```

- ▶ Together `group_by()` and `summarize()` are one the most popular tools in **dplyr**.

Combining Multiple Operations with the Pipe

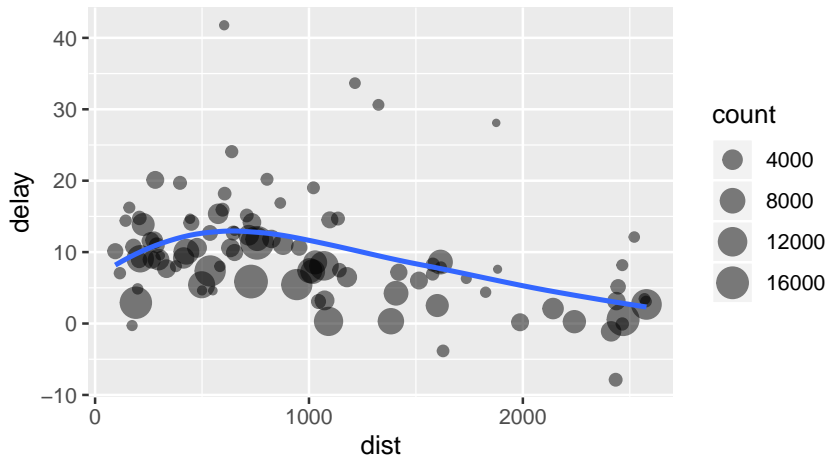
- ▶ You want to explore the relationship between the distance and average delay for each location.

```
by_dest <- group_by(flights, dest)
delay <- summarize(by_dest,
                   count = n(),
                   dist = mean(distance, na.rm = TRUE),
                   delay = mean(arr_delay, na.rm = TRUE)
                  )
delay <- filter(delay, count > 20, dest != "HNL")
```

Combining Multiple Operations with the Pipe

```
ggplot(data = delay, mapping = aes(x = dist, y = delay)) +  
  geom_point(aes(size = count), alpha = 1/2) +  
  geom_smooth(se = FALSE)
```

`geom_smooth()` using method = 'loess' and formula 'y ~ x'



Combining Multiple Operations with the Pipe

- Equivalent expression with the pipe `%>%`.

```
delays <- flights %>%  
  group_by(dest) %>%  
  summarize(  
    count = n(),  
    dist = mean(distance, na.rm = TRUE),  
    delay = mean(arr_delay, na.rm = TRUE)  
  ) %>%  
  filter(count > 20, dest != "HNL")
```

Missing Values

```
not_cancelled <- flights %>%  
  filter(!is.na(dep_delay), !is.na(arr_delay))  
not_cancelled %>%  
  group_by(year, month, day) %>%  
  summarize(mean = mean(dep_delay))
```

```
## # A tibble: 365 x 4  
## # Groups:   year, month [12]  
##   year month   day mean  
##   <int> <int> <int> <dbl>  
## 1  2013     1     1  11.4  
## 2  2013     1     2  13.7  
## 3  2013     1     3  10.9  
## 4  2013     1     4   8.97  
## 5  2013     1     5   5.73  
## 6  2013     1     6   7.15  
## 7  2013     1     7   5.42  
## 8  2013     1     8   2.56  
## 9  2013     1     9   2.30  
## 10 2013     1    10   2.84  
## # ... with 355 more rows
```


Useful Summary Functions (location)

```
not_cancelled %>%
  group_by(year, month, day) %>%
  summarize(
    # averaged delay:
    ave_delay1 = mean(arr_delay),
    # averaged positive delay:
    ave_delay2 = mean(arr_delay[arr_delay > 0])
  )
```

```
## # A tibble: 365 x 5
## # Groups:   year, month [12]
##   year month   day ave_delay1 ave_delay2
##   <int> <int> <int>      <dbl>      <dbl>
## 1  2013     1     1      12.7       32.5
## 2  2013     1     2      12.7       32.0
## 3  2013     1     3       5.73      27.7
## 4  2013     1     4      -1.93      28.3
## 5  2013     1     5      -1.53      22.6
## 6  2013     1     6       4.24      24.4
## 7  2013     1     7      -4.95      27.8
## 8  2013     1     8      -3.23      20.8
## 9  2013     1     9      -0.264     25.6
## 10 2013     1    10      -5.90      27.3
```

Useful Summary Functions (dispersion)

```
not_cancelled %>%  
  group_by(dest) %>%  
  summarize(distance_sd = sd(distance)) %>%  
  arrange(desc(distance_sd))
```

```
## # A tibble: 104 x 2  
##   dest distance_sd  
##   <chr>         <dbl>  
## 1 EGE          10.5  
## 2 SAN          10.4  
## 3 SFO          10.2  
## 4 HNL          10.0  
## 5 SEA           9.98  
## 6 LAS           9.91  
## 7 PDX           9.87  
## 8 PHX           9.86  
## 9 LAX           9.66  
## 10 IND          9.46  
## # ... with 94 more rows
```

Useful Summary Functions (others)

- ▶ rank: `min(x)`, `quantile(x, 0.25)`, `max(x)`
- ▶ position: `first(x)`, `nth(x, 2)`, `last(x)`
- ▶ counts: `n()`, `n_distinct()`
- ▶ Others: `sum(x > 10)`, `mean(y == 0)`

Grouping by Multiple Variables

```
daily <- group_by(flights, year, month, day)
per_day <- summarize(daily, flights = n())
per_day
```

```
## # A tibble: 365 x 4
## # Groups:   year, month [12]
##   year month   day flights
##   <int> <int> <int>    <int>
## 1  2013     1     1     842
## 2  2013     1     2     943
## 3  2013     1     3     914
## 4  2013     1     4     915
## 5  2013     1     5     720
## 6  2013     1     6     832
## 7  2013     1     7     933
## 8  2013     1     8     899
## 9  2013     1     9     902
## 10 2013     1    10     932
## # ... with 355 more rows
```

Ungrouping

```
daily %>%  
  ungroup() %>%  
  summarize(flights = n())
```

```
## # A tibble: 1 x 1  
##   flights  
##   <int>  
## 1  336776
```

Grouped Mutates (and Filters)

- Find the worst members of each group

```
flights_sml %>%  
  group_by(year, month, day) %>%  
  filter(rank(desc(arr_delay)) < 10)
```

```
## # A tibble: 3,306 x 7  
## # Groups:   year, month, day [365]  
##   year month   day dep_delay arr_delay distance air_time  
##   <int> <int> <int>     <dbl>     <dbl>     <dbl>     <dbl>  
## 1  2013     1     1       853       851       184        41  
## 2  2013     1     1       290       338      1134       213  
## 3  2013     1     1       260       263       266        46  
## 4  2013     1     1       157       174       213        60  
## 5  2013     1     1       216       222       708       121  
## 6  2013     1     1       255       250       589       115  
## 7  2013     1     1       285       246      1085       146  
## 8  2013     1     1       192       191       199        44  
## 9  2013     1     1       379       456      1092       222  
## 10 2013     1     2       224       207       550        94  
## # ... with 3,296 more rows
```

Grouped Mutates (and Filters)

- Find all groups bigger than a threshold.

```
flights %>%  
  group_by(dest) %>%  
  filter(n() > 365)
```

```
## # A tibble: 332,577 x 19  
## # Groups:   dest [77]  
##   year month   day dep_time sched_dep_time dep_delay arr_time  
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>  
## 1  2013     1     1     517           515           2     830  
## 2  2013     1     1     533           529           4     850  
## 3  2013     1     1     542           540           2     923  
## 4  2013     1     1     544           545          -1    1004  
## 5  2013     1     1     554           600          -6     812  
## 6  2013     1     1     554           558          -4     740  
## 7  2013     1     1     555           600          -5     913  
## 8  2013     1     1     557           600          -3     709  
## 9  2013     1     1     557           600          -3     838  
## 10 2013     1     1     558           600          -2     753  
## # ... with 332,567 more rows, and 12 more variables: sched_arr_time  
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,  
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <  
## #   minute <dbl>, time_hour <dtm>
```

Reference

- ▶ Wickham, H. and Grolemund, G. (2017) R for Data Science, O'reilly Media Inc., Chapter 4.