

# ST720 Data Science

## Handling Strings

Seung Jun Shin (sjshin@krea.ac.kr)

Department of Statistics, Korea University

# Introduction

- ▶ Base **R** contains many functions to work with strings but not convenient at all.
- ▶ `stringr` package is more intuitive.

# String basics

## ► String Length

```
str_length(c("a", "R for data science", NA))
```

```
## [1]  1 18 NA
```

## ► Combining Strings

```
str_c("x", "y")
```

```
## [1] "xy"
```

```
str_c("x", "y", sep=",")
```

```
## [1] "x,y"
```

## ► Handling Missing

```
x <- c('abc', NA)  
str_c("|-", x, "-|")
```

```
## [1] "|-abc-|" NA
```

```
str_c("|-", str_replace_na(x), "-|")
```

```
## [1] "|-abc-|" "|-NA-|"
```

# Combining Strings

- ▶ More complicating.

```
name <- 'Hadley'  
time_of_day <- 'morning'  
birthday <- FALSE  
str_c('Good ', time_of_day, " ", name,  
      if(birthday) 'and HAPPY BIRTHDAY', '.')
```

```
## [1] "Good morning Hadley."
```

- ▶ Collapse a vector into a single string

```
str_c(c("x", "y", "z"), collapse= ",")
```

```
## [1] "x,y,z"
```

# Subsetting Strings

```
x <- c("Apple", "Banana", "Pear")  
str_sub(x, 1, 3)
```

```
## [1] "App" "Ban" "Pea"
```

```
str_sub(x, -3, -1)
```

```
## [1] "ple" "ana" "ear"
```

```
str_sub("a", 1, 5)
```

```
## [1] "a"
```

```
str_sub(x, 1, 1) <- str_to_lower(str_sub(x,1,1))  
x
```

```
## [1] "apple" "banana" "pear"
```

# Locales

```
dog <- "The quick brown dog"  
str_to_upper(dog)  
str_to_lower(dog)  
str_to_title(dog)
```

```
## [1] "THE QUICK BROWN DOG"
```

```
## [1] "the quick brown dog"
```

```
## [1] "The Quick Brown Dog"
```

- ▶ Turkish has two i's : with and without a dot, and it has a different rule for capitalizing them :

```
str_to_upper(c("i", "ı"))  
str_to_upper(c("i", "ı"), locale = "tr")
```

```
## [1] "I" "I"
```

```
## [1] "İ" "I"
```

# Locales

```
x <- c("apple", "eggplant", "banana")  
str_sort(x, locale = 'en') # English
```

```
## [1] "apple"      "banana"     "eggplant"
```

```
str_sort(x, locale = 'haw') # Hawaiian
```

```
## [1] "apple"      "eggplant"   "banana"
```

- ▶ The locale is specified as an ISO 639 language code, which is a two or three-letter abbreviation.

# Regulars Expressions



# Regexps

**Regexps** are a very terse language that allow you to describe patterns in strings.

메타문자	기능	설명
.	문자	1개의 문자와 일치한다. 단일행 모드에서는 <b>새줄 문자</b> 를 제외한다.
[ ]	문자 클래스	"["과 "]" 사이의 문자 중 하나를 선택한다. " "를 여러 개 쓴 것과 같은 의미이다. 예를 들면 [abc]d는 ad, bd, cd를 뜻한다. 또한, "-" 기호와 함께 쓰면 범위를 지정할 수 있다. "[a-z]"는 a부터 z까지 중 하나, "[1-9]"는 1부터 9까지 중의 하나를 의미한다.
[^ ]	부정	문자 클래스 안의 문자를 제외한 나머지를 선택한다. 예를 들면 [^abc]d는 ad, bd, cd는 포함하지 않고 ed, fd 등을 포함한다. ["a-z]"는 알파벳 소문자로 시작하지 않는 모든 문자를 의미한다.
^	처음	문자열이나 행의 처음을 의미한다.
\$	끝	문자열이나 행의 끝을 의미한다.
( )	하위식	여러 식을 하나로 묶을 수 있다. "abc adc"와 "a(b d)c"는 같은 의미를 가진다.
\n	일치하는 n번째 패턴	일치하는 패턴들 중 n번째를 선택하며, 여기에서 n은 1에서 9 중 하나가 올 수 있다.
*	0회 이상	0개 이상의 문자를 포함한다. "a*b"는 "b", "ab", "aab", "aaaab"를 포함한다.
{m, n}	m회 이상 n회 이하	"a{1,3}b"는 "ab", "aab", "aaaab"를 포함하지만, "b"나 "aaaaab"는 포함하지 않는다.

Figure 1: regexp1

# Basic Matches

- ▶ The simplest patterns match exact strings:

```
x <- c("apple", "banana", "pear")  
str_view(x, "an")
```

apple

bananana

pear

- ▶ The next step up in complexity is ., which matches any character :

```
str_view(x, ".a.")
```

apple

bananana

pear

# Basic Matches

- ▶ To create the regular expression `\.` we need the string `"\\."`

```
dot <- "\\."
```

- ▶ But the expression itself only contains one:

```
writeLines(dot)
```

```
## \.
```

- ▶ And this tells R to look for an explicit `.`

```
str_view(c("abc", "a.c", "bef"), "a\\.c")
```

abc

a.c

bef

# Anchors

- ▶ By default, regular expressions will match any part of a string. It's often useful to *anchor* the regular expression so that it matches from the start or end of the string.
  - ▶ `^` to match the start of the string.
  - ▶ `$` to match the end of the string.

```
x <- c("apple", "banana", "pear")  
str_view(x, "^a")
```

apple

banana

pear

```
str_view(x, "a$")
```

apple

banana

## Anchors

- ▶ To force a regular expression to only match a complete string, anchor it with both `^` and `$` :

```
x <- c("apple pie", "apple", "apple cake")  
str_view(x, "apple")
```

apple pie

apple

apple cake

```
str_view(x, "^apple$")
```

apple pie

apple

apple cake

# Character Classes and Alternatives

- ▶ There are a number of special patterns that match more than one character.
  - ▶ `\d`: matches any digit.
  - ▶ `\s`: matches any whitespace (e.g. space, tab, newline).
  - ▶ `[abc]`: matches a, b, or c.
  - ▶ `[^abc]`: matches anything except a, b, or c.
- ▶ Remember, to create a regular expression containing `\d` or `\s`, you'll need to escape the `\` for the string, so you'll type `\\d` or `\\s`.
- ▶ A character class containing a single character is a nice alternative to backslash escapes when you want to include a single metacharacter in a regex.

# Character Classes and Alternatives

```
str_view(c("abc", "a.c", "a*c", "a c"), "a[.]c")
```

abc

a.c

a\*c

a c

# Character Classes and Alternatives

```
str_view(c("abc", "a.c", "a*c", "a c"), "[*]c")
```

abc

a.c

a\*c

a c



## Character Classes and Alternatives

```
str_view(c("abc", "a.c", "a*c", "a c"), "a[ ]")
```

abc

a.c

a\*c

a c

# Character Classes and Alternatives

- ▶ You can use *alternation* to pick between one or more alternative patterns.
- ▶ For example, `abc|d..f` will match either “abc”, or “deaf”.

```
str_view(c("grey", "gray"), "gr(e|a)y")
```

grey

gray

# Repetition

- ▶ ? : 0 or 1
- ▶ + : 1 or more
- ▶ \* : 0 or more

```
x <- "1888 is the longest year in Roman numerals: MDCCCLXXXVIII"  
str_view(x, "CC?")
```

1888 is the longest year in Roman numerals: MDCCCLXXXVIII

```
str_view(x, "CC+")
```

1888 is the longest year in Roman numerals: MDCCCLXXXVIII

```
str_view(x, 'C[LX]+')
```

1888 is the longest year in Roman numerals: MDCCCLXXXVIII

# Repetition

- ▶ We can also specify the number of matches precisely :
  - ▶  $\{n\}$  : exactly  $n$
  - ▶  $\{n, \}$  :  $n$  or more
  - ▶  $\{, m\}$  : at most  $m$
  - ▶  $\{n, m\}$  : between  $n$  and  $m$

## Repetition

```
str_view(x, "C{2}")
```

1888 is the longest year in Roman numerals: MDCCCLXXXVIII

```
str_view(x, "C{2,}")
```

1888 is the longest year in Roman numerals: MDCCCLXXXVIII

```
str_view(x, "C{2,3}")
```

1888 is the longest year in Roman numerals: MDCCCLXXXVIII

```
str_view(x, 'C{2,3}?')
```

1888 is the longest year in Roman numerals: MDCCCLXXXVIII

```
str_view(x, 'C[LX]+?')
```

# Grouping and Backreferences

- ▶ Parentheses also create a numbered capturing group (number 1, 2 etc.).
- ▶ A capturing group stores the part of the string matched by the part of the regular expression inside the parentheses.
- ▶ You can refer to the same text as previously matched by a capturing group with backreferences, like `\1`, `\2` etc.
- ▶ For example, the following regular expression finds all fruits that have a repeated pair of letters.

## Grouping and Backreferences

```
str_view(fruit, "(...)\1", match = TRUE)
```

banana

coconut

cucumber

jujube

papaya

sala berry

# Detect Matches

- ▶ To determine if a character vector matches a pattern, use `str_detect`. It returns a logical vector the same length as the input:

```
x <- c("apple", "banana", "pear")  
str_detect(x, "e")
```

```
## [1] TRUE FALSE TRUE
```

- ▶ How many common words start with t?

```
sum(str_detect(words, "^t"))
```

```
## [1] 65
```

- ▶ What proportion of common words end with a vowel?

```
mean(str_detect(words, '[aeiou]$'))
```

```
## [1] 0.2765306
```



# Detect Matches

- For example, here are two ways to find all words that don't contain any vowels:

```
no_vowels1 <- !str_detect(words, '[aeiou]')  
no_vowels2 <- str_detect(words, '^[^aeiou]+$')  
identical(no_vowels1, no_vowels2)
```

```
## [1] TRUE
```

# Detect Matches

- Typically, however, your strings will be one column of a data frame, and you'll want to use `filter` instead:

```
df <- tibble(  
  word <- words,  
  i = seq_along(words)  
)  
df %>%  
  filter(str_detect(words, "x$"))
```

```
## # A tibble: 4 x 2  
##   `word <- words`      i  
##   <chr>             <int>  
## 1 box               108  
## 2 sex               747  
## 3 six               772  
## 4 tax               841
```

# Detect Matches

- ▶ A variation on `str_detect()` is `str_count()`: rather than a simple yes or no, it tells you how many matches there are in a string:

```
x <- c("apple", "banana", "pear")  
str_count(x, "a")
```

```
## [1] 1 3 1
```

- ▶ On average, how many vowels per word?

```
mean(str_count(words, "[aeiou]"))
```

```
## [1] 1.991837
```

## Extract Matches

- ▶ To extract the actual text of a match, use `str_extract()`.
- ▶ To show that off, we're going to need a more complicated example. I'm going to use the **Harvard sentences**, which were designed to test VOIP systems, but are also useful for practicing regexes.

```
length(sentences)
```

```
## [1] 720
```

```
head(sentences)
```

```
## [1] "The birch canoe slid on the smooth planks."  
## [2] "Glue the sheet to the dark blue background."  
## [3] "It's easy to tell the depth of a well."  
## [4] "These days a chicken leg is a rare dish."  
## [5] "Rice is often served in round bowls."  
## [6] "The juice of lemons makes fine punch."
```

# Extract Matches

```
colors <- c("red", "orange", "yellow", "green", "blue", "purple")
color_match <- str_c(colors, collapse = "|")
color_match
```

```
## [1] "red|orange|yellow|green|blue|purple"
```

```
has_color <- str_subset(sentences, color_match)
matches <- str_extract(has_color, color_match)
head(matches)
```

```
## [1] "blue" "blue" "red" "red" "red" "blue"
```

- Note that `str_extract()` only extracts the **first** match.

## Extract Matches

- To get all matches, use `str_extract_all()`.

```
more <- sentences[str_count(sentences, color_match) > 1]  
str_extract_all(more, color_match)
```

```
## [[1]]  
## [1] "blue" "red"  
##  
## [[2]]  
## [1] "green" "red"  
##  
## [[3]]  
## [1] "orange" "red"
```

# Grouped Matches

- ▶ Imagine we want to extract nouns from the sentences. As a heuristic, we'll look for any word that comes after "a" or "the".
- ▶ Defining a "word" in a regular expression is a little tricky, so here I use a simple approximation a sequence of at least one character that isn't a space:

```
noun <- "(a|the) ([^ ]+)"  
has_noun <- sentences %>%  
  str_subset(noun)%>%  
  head(10)  
has_noun %>%  
  str_extract(noun)
```

```
## [1] "the smooth" "the sheet" "the depth" "a chicken" "the  
## [6] "the sun" "the huge" "the ball" "the woman" "a h
```

## Grouped Matches

- ▶ `str_extract()` gives us the complete match; `str_match()` gives each individual component.

```
has_noun %>%  
  str_match(noun)
```

```
##      [,1]      [,2] [,3]  
## [1,] "the smooth" "the" "smooth"  
## [2,] "the sheet"  "the" "sheet"  
## [3,] "the depth"  "the" "depth"  
## [4,] "a chicken"  "a"   "chicken"  
## [5,] "the parked" "the" "parked"  
## [6,] "the sun"     "the" "sun"  
## [7,] "the huge"   "the" "huge"  
## [8,] "the ball"   "the" "ball"  
## [9,] "the woman"  "the" "woman"  
## [10,] "a helps"   "a"   "helps"
```



## Grouped Matches

- ▶ If your data is in tibble, it's often easier to use `tidyr::extract()`. It works like `str_match()` but requires you to name the matches, which are then placed in new columns.

```
tibble(sentence = sentences) %>%  
  tidyr::extract(  
    sentence, c("article", "noun"), "(a|the) ([^ ]+)",  
    remove = FALSE  
  )
```

```
## # A tibble: 720 x 3
```

##	sentence	article	noun
##	<chr>	<chr>	<chr>
## 1	The birch canoe slid on the smooth planks.	the	smooth
## 2	Glue the sheet to the dark blue background.	the	sheet
## 3	It's easy to tell the depth of a well.	the	depth
## 4	These days a chicken leg is a rare dish.	a	chicke
## 5	Rice is often served in round bowls.	<NA>	<NA>
## 6	The juice of lemons makes fine punch.	<NA>	<NA>
## 7	The box was thrown beside the parked truck.	the	parked
## 8	The hogs were fed chopped corn and garbage.	<NA>	<NA>
## 9	Four hours of steady work faced us.	<NA>	<NA>

# Replacing Matches

- ▶ `str_replace()` and `str_replace_all()` allow you to replace matches with new strings. The simplest use is to replace a pattern with a fixed string:

```
x <- c("apple", "pear", "banana")  
str_replace(x, "[aeiou]", "-")
```

```
## [1] "-pple" "p-ar"  "b-nana"
```

```
str_replace_all(x, "[aeiou]", "-")
```

```
## [1] "-ppl-" "p--r"  "b-n-n-"
```

# Splitting

Use `str_split()` to split a string up into pieces. For example, we could split sentences into words:

```
sentences %>%  
  head(5) %>%  
  str_split(" ")
```

```
## [[1]]  
## [1] "The"      "birch"    "canoe"    "slid"     "on"       "the"      "smooth"  
## [8] "planks."  
##  
## [[2]]  
## [1] "Glue"      "the"      "sheet"    "to"       "the"  
## [6] "dark"      "blue"     "background."  
##  
## [[3]]  
## [1] "It's"     "easy"     "to"       "tell"     "the"      "depth"   "of"      "a"        "well."  
##  
## [[4]]  
## [1] "These"    "days"    "a"        "chicken"  "leg"      "is"      "a"  
## [8] "rare"     "dish."  
##  
## [[5]]  
## [1] "Rice"     "is"       "often"    "served"   "in"       "round"   "bowls."
```

## Splitting

- ▶ We can also request a maximum number of pieces:

```
fields <- c("Name: Hadley", "Country: NZ", "Age: 35")
fields %>% str_split(":", n=2, simplify = TRUE)
```

```
##      [,1]      [,2]
## [1,] "Name"    "Hadley"
## [2,] "Country" "NZ"
## [3,] "Age"     "35"
```

- ▶ Instead of splitting up strings by patterns, you can also split up by character, line, sentence, and word boundary():

```
x <- "This is a sentence. This is another sentence."
str_view_all(x, boundary("word"))
```

This is a sentence. This is another sentence.

```
str_split(x, " ")[[1]]
```

```
## [1] "This"      "is"         "a"          "sentence." ""
## [7] "is"        "another"    "sentence."
```

## Other Types of Pattern

- ▶ When you use a pattern that's string, it's automatically wrapped into a call to `regex()`:
- ▶ `ignore_case = TRUE` allows characters to match either their uppercase or lowercase forms. This always uses the current locale:

```
bananas <- c("banana", "Banana", "BANANA")  
str_view(bananas, "banana")
```

banana

Banana

BANANA

```
str_view(bananas, regex("banana", ignore_case = TRUE))
```

banana

Banana

BANANA

## Other Types of Pattern

- ▶ `multiline = TRUE` allows `^` and `$` to match the start and end of each line rather than the start and end of the complete string.

```
x <- "Line 1\nLine 2\nLine 3"  
str_extract_all(x, "^Line")[[1]]
```

```
## [1] "Line"
```

```
str_extract_all(x, regex("^Line", multiline = TRUE))[[1]]
```

```
## [1] "Line" "Line" "Line"
```

## Other Types of Pattern

- ▶ `comments = TRUE` allows you to use comments and white space to make complex regular expressions more understandable. Spaces are ignored, as is everything after `#`. To match a literal space, you'll need to escape it : `//`

```
phone <- regex("
    \\(?:      # optional opening parens
    (\\d{3})   # area code
    [- ]      # optional closing parens, dash, or space
    (\\d{3})   # another three numbers
    [ -]      # optional space or dash
    (\\d{3})   # three more numbers
    ", comments = TRUE)
str_match('123-456-7890', phone)
```

```
##      [,1]      [,2]  [,3]  [,4]
## [1,] "123-456-789" "123" "456" "789"
```

## stringi

- ▶ **stringr** is built on top of the stringi package. **stringr** is useful when you're learning because it exposes a minimal set of functions, which have been carefully picked to handle the most common string manipulation functions.
- ▶ **stringi**, on the other hand, is designed to be comprehensive. It contains almost every function you might ever need: stringi has 234 functions to stringr's 42. The main difference is the prefix: `str_` versus `stri_`.