

ST720 Data Science

Sentiment analysis with tidy data

Seung Jun Shin (sjshin@krea.ac.kr)

Department of Statistics, Korea University

Introduction

- ▶ **Word frequency** allowed us to analyze which words are used most frequently in documents and to compare documents.
- ▶ When human readers approach a text, we use our understanding of the emotional intent of words to infer whether a section of text is positive or negative.
- ▶ We can use the tools of text mining to approach the emotional content of text programmatically.
- ▶ One way to analyze the sentiment of a text is to consider the text as a combination of its individual words and the sentiment content of the whole text as the sum of the sentiment content of the individual words.

sentiments Dataset

- ▶ The tidytext package contains several sentiment lexicons
 - ▶ AFINN: assigns words with a score that runs between -5 and 5, with negative scores indicating negative sentiment and positive scores indicating positive sentiment.
 - ▶ `bing` : categorizes words in a binary fashion into positive and negative categories.
 - ▶ `nrc`: categorizes words in a binary fashion (“yes”/“no”) into categories of positive, negative, anger, anticipation, disgust, fear, joy, sadness, surprise, and trust.
- ▶ Contain many English words and the words are assigned scores for positive/negative sentiment, and also possibly emotions like joy, anger, sadness, and so forth.

sentiments Dataset

- All of this information is tabulated in the sentiments dataset, and tidytext provides a function `get_sentiments()` to get specific sentiment lexicons.

```
library(tidytext)
print(get_sentiments("afinn"), n = 5)
```

```
## # A tibble: 2,477 x 2
##   word      value
##   <chr>    <dbl>
## 1 abandon      -2
## 2 abandoned    -2
## 3 abandons     -2
## 4 abducted     -2
## 5 abduction    -2
## # ... with 2,472 more rows
```

sentiments Dataset

```
get_sentiments("bing")
```

```
## # A tibble: 6,786 x 2
```

```
##   word      sentiment
```

```
##   <chr>      <chr>
```

```
## 1 2-faces    negative
```

```
## 2 abnormal  negative
```

```
## 3 abolish   negative
```

```
## 4 abominable negative
```

```
## 5 abominably negative
```

```
## 6 abominate  negative
```

```
## 7 abomination negative
```

```
## 8 abort      negative
```

```
## 9 aborted    negative
```

```
## 10 abortions  negative
```

```
## # ... with 6,776 more rows
```

sentiments Dataset

```
get_sentiments("nrc")
```

```
## # A tibble: 13,901 x 2
##   word      sentiment
##   <chr>    <chr>
## 1 abacus    trust
## 2 abandon   fear
## 3 abandon   negative
## 4 abandon   sadness
## 5 abandoned anger
## 6 abandoned fear
## 7 abandoned negative
## 8 abandoned sadness
## 9 abandonment anger
## 10 abandonment fear
## # ... with 13,891 more rows
```

Dictionary-based Sentiment analysis

- ▶ Dictionary-based methods like the ones we are discussing find the total sentiment of a piece of text by adding up the individual sentiment scores for each word in the text.

Sentiment analysis with inner join

```
library(janeaustenr)
library(dplyr)
library(stringr)

tidy_books <- austen_books() %>%
  group_by(book) %>%
  mutate(linenumber = row_number(),
         chapter = cumsum(
           str_detect(text,
                     regex("^chapter [\\divxlc]",
                           ignore_case = TRUE)))) %>%
  ungroup() %>%
  unnest_tokens(word, text)
```


Sentiment analysis with inner join

```
nrc_joy <- get_sentiments("nrc") %>%  
  filter(sentiment == "joy")  
tidy_books %>%  
  filter(book == "Emma") %>%  
  inner_join(nrc_joy) %>%  
  count(word, sort = TRUE) %>%  
  print(n = 5)
```

```
## Joining, by = "word"
```

```
## # A tibble: 303 x 2
```

```
##   word      n
```

```
##   <chr>  <int>
```

```
## 1 good    359
```

```
## 2 young   192
```

```
## 3 friend  166
```

```
## 4 hope    143
```

```
## 5 happy   125
```

```
## # ... with 298 more rows
```

Sentiment analysis with inner join

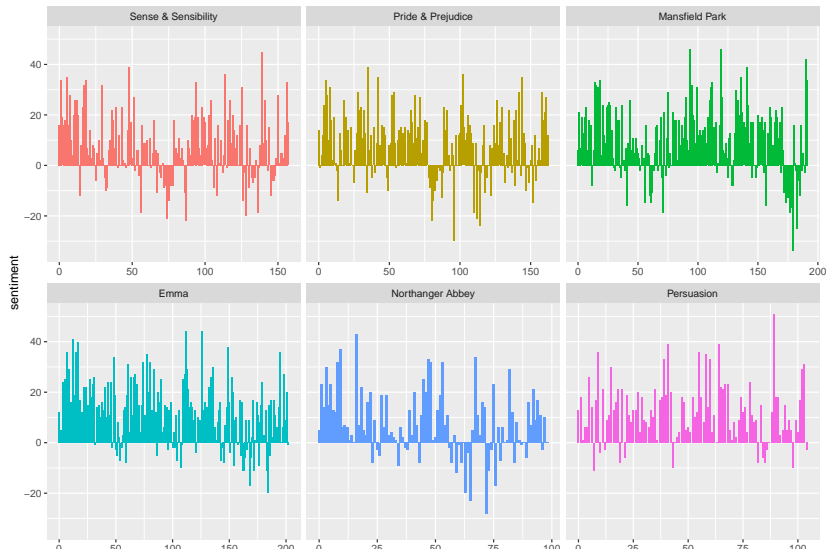
```
jane_austen_sentiment <- tidy_books %>%  
  inner_join(get_sentiments("bing")) %>%  
  count(book, index = linenumber %/% 80, sentiment) %>%  
  spread(sentiment, n, fill = 0) %>%  
  mutate(sentiment = positive - negative)  
print(jane_austen_sentiment, n = 5)
```

```
## # A tibble: 920 x 5
```

##	book	index	negative	positive	sentiment
##	<fct>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	Sense & Sensibility	0	16	32	16
## 2	Sense & Sensibility	1	19	53	34
## 3	Sense & Sensibility	2	12	31	19
## 4	Sense & Sensibility	3	15	31	16
## 5	Sense & Sensibility	4	16	34	18
## #	... with 915 more rows				

Sentiment analysis with inner join

```
ggplot(jane_austen_sentiment, aes(index, sentiment, fill = book))  
  geom_col(show.legend = FALSE) +  
  facet_wrap(~book, ncol = 3, scales = "free_x")
```



Most common positive and negative words

- Analyze word counts that contribute to each sentiment.

```
bing_word_counts <- tidy_books %>%  
  inner_join(get_sentiments("bing")) %>%  
  count(word, sentiment, sort = TRUE) %>%  
  ungroup()
```

```
## Joining, by = "word"
```

```
print(bing_word_counts, n = 5)
```

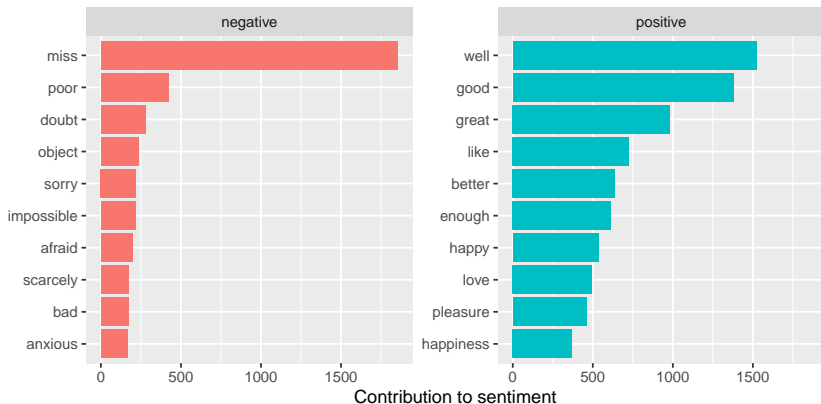
```
## # A tibble: 2,585 x 3  
##   word    sentiment      n  
##   <chr> <chr>      <int>  
## 1 miss   negative    1855  
## 2 well   positive    1523  
## 3 good   positive    1380  
## 4 great  positive     981  
## 5 like   positive     725  
## # ... with 2,580 more rows
```

Most common positive and negative words

- ▶ Let's Visualize it

```
bing_word_counts %>%  
  group_by(sentiment) %>%  
  top_n(10) %>%  
  ungroup() %>%  
  mutate(word = reorder(word, n)) %>%  
  ggplot(aes(word, n, fill = sentiment)) +  
  geom_col(show.legend = FALSE) +  
  facet_wrap(~sentiment, scales = "free_y") +  
  labs(y = "Contribution to sentiment",  
       x = NULL) +  
  coord_flip()
```

Most common positive and negative words



Stop Words Editing

- ▶ **miss** is coded as negative but it is used as a title for young.
- ▶ Add “miss” to a custom stop-words list using `bind_rows()`

```
custom_stop_words <- bind_rows(tibble(word = c("miss"),  
                                       lexicon = c("custom")),  
                               stop_words)  
print(custom_stop_words, n = 5)
```

```
## # A tibble: 1,150 x 2  
##   word  lexicon  
##   <chr> <chr>  
## 1 miss  custom  
## 2 a     SMART  
## 3 a's   SMART  
## 4 able  SMART  
## 5 about SMART  
## # ... with 1,145 more rows
```

Wordclouds

- Tidy format is useful for wordclouds in wordcloud package.

```
tidy_books %>%  
  anti_join(stop_words) %>%  
  count(word) %>%  
  with(wordcloud(word, n, max.words = 100))
```


Wordclouds



Wordclouds: Comparison

```
library(reshape2)
tidy_books %>%
  inner_join(get_sentiments("bing")) %>%
  count(word, sentiment, sort = TRUE) %>%
  acast(word ~ sentiment, value.var = "n", fill = 0) %>%
  comparison.cloud(colors = c("gray20", "gray80"),
                   max.words = 100)
```

- Check acast function in reshape2 package.

Wordclouds: Comparison

negative



positive

Looking at units beyond just words

- ▶ Sometimes it is useful or necessary to look beyond single words.
- ▶ For example, one can tokenize text in a sentence level:

```
PandP_sentences <- tibble(text = prideprejudice) %>%  
  unnest_tokens(sentence, text, token = "sentences")  
PandP_sentences$sentence[2]
```

```
## [1] "however little known the feelings or views of such a man
```

Looking at units beyond just words

- ▶ Token can be expressed as regex.

```
austen_chapters <- austen_books() %>%  
  group_by(book) %>%  
  unnest_tokens(chapter, text, token = "regex",  
                pattern = "Chapter|CHAPTER [\\dIVXLC]") %>%  
  ungroup()  
  
austen_chapters %>% group_by(book) %>%  
  summarise(chapters = n()) %>% print(n = 3)
```

```
## # A tibble: 6 x 2  
##   book          chapters  
##   <fct>         <int>  
## 1 Sense & Sensibility      51  
## 2 Pride & Prejudice       62  
## 3 Mansfield Park          49  
## # ... with 3 more rows
```

Looking at units beyond just words

- Which chapter is most negative in each of Jane Austen's novels?

```
bingnegative <- get_sentiments("bing") %>%  
  filter(sentiment == "negative")
```

```
wordcounts <- tidy_books %>%  
  group_by(book, chapter) %>%  
  summarize(words = n())  
print(wordcounts, n = 5)
```

```
## # A tibble: 275 x 3  
## # Groups:   book [6]  
##   book                chapter words  
##   <fct>                <int> <int>  
## 1 Sense & Sensibility      0      7  
## 2 Sense & Sensibility      1 1571  
## 3 Sense & Sensibility      2 1970  
## 4 Sense & Sensibility      3 1538  
## 5 Sense & Sensibility      4 1952  
## # ... with 270 more rows
```

Looking at units beyond just words

```
tidy_books %>%  
  inner_join(bingnegative) %>%  
  group_by(book, chapter) %>%  
  summarize(negativewords = n()) %>%  
  left_join(wordcounts, by = c("book", "chapter")) %>%  
  mutate(ratio = negativewords/words) %>%  
  filter(chapter != 0) %>%  
  top_n(1) %>%  
  ungroup()
```

```
## # A tibble: 6 x 5  
##   book                chapter negativewords words  ratio  
##   <fct>              <int>          <int> <int> <dbl>  
## 1 Sense & Sensibility    43            161  3405 0.0473  
## 2 Pride & Prejudice      34            111  2104 0.0528  
## 3 Mansfield Park        46            173  3685 0.0469  
## 4 Emma                   15            151  3340 0.0452  
## 5 Northanger Abbey      21            149  2982 0.0500  
## 6 Persuasion             4             62  1807 0.0343
```

Summary

- ▶ Sentiment analysis provides
 - ▶ a way to understand the attitudes and opinions expressed in texts.
 - ▶ how a narrative arc changes throughout its course or what words with emotional and opinion content are important for a particular text.
- ▶ With tidy text data, sentiment analysis can be implemented as an `inner_join`.