

# ST720 Data Science

tf-idf

Seung Jun Shin (sjshin@krea.ac.kr)

Department of Statistics, Korea University

# Introduction

- ▶ How to quantify what a document is about?
- ▶ Can we do this by looking at the words?
- ▶ One measure of how important a word is its **term frequency (tf)**.
- ▶ There are words in a document occur many times but not important.
- ▶ Some eliminated words are more important in some documents than others.
- ▶ One way is to use **a term's inverse document frequency (idf)**.

$$\text{idf}(\text{term}) = \log \left( \frac{n_{\text{documents}}}{n_{\text{documents containing term}}} \right)$$

- ▶ **tf-idf** (the two quantities multiplied together) is the frequency of a term adjusted for how rarely it is used.

# Term frequency in Jane Austen's novels

- What are the most commonly used words in Jane Austen's novels?

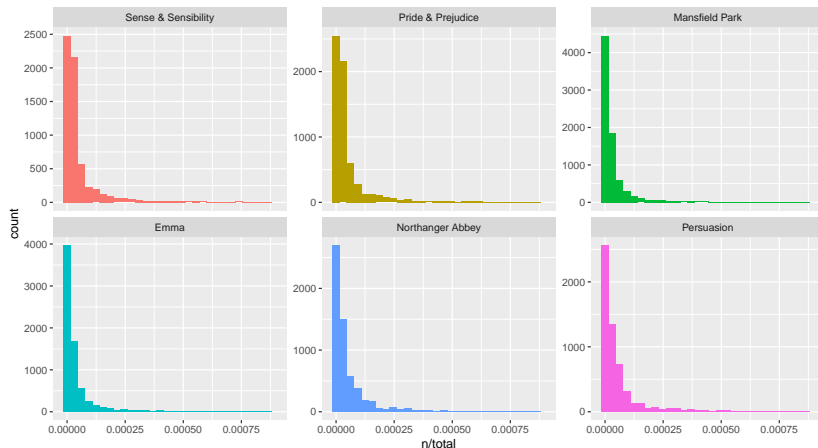
```
book_words <- austen_books() %>%  
  unnest_tokens(word, text) %>%  
  count(book, word, sort = TRUE)  
total_words <- book_words %>%  
  group_by(book) %>%  
  summarize(total = sum(n))  
book_words <- left_join(book_words, total_words)  
print(book_words, n = 4)
```

```
## # A tibble: 40,379 x 4  
##   book          word      n  total  
##   <fct>         <chr> <int> <int>  
## 1 Mansfield Park the      6206 160460  
## 2 Mansfield Park to       5475 160460  
## 3 Mansfield Park and      5438 160460  
## 4 Emma          to       5239 160996  
## # ... with 4.038e+04 more rows
```

# Term frequency in Jane Austen's novels

- Distribution of  $n/\text{total}$  for each novel.

```
ggplot(book_words, aes(n/total, fill = book)) +  
  geom_histogram(show.legend = FALSE) +  
  xlim(NA, 0.0009) +  
  facet_wrap(~book, ncol = 3, scales = "free_y")
```



# Zipf's Law

- ▶ Distributions with a long tail is typical in language.
- ▶ Zipf's Law states

$$tf \propto \frac{1}{\text{rank}}$$


```
freq_by_rank <- book_words %>%  
  group_by(book) %>%  
  mutate(rank = row_number(),  
         `term frequency` = n/total)  
print(freq_by_rank, n = 3)
```

```
## # A tibble: 40,379 x 6  
## # Groups:   book [6]  
##   book          word      n  total  rank `term frequency`  
##   <fct>         <chr> <int>  <int> <int>          <dbl>  
## 1 Mansfield Park the      6206 160460     1      0.0387  
## 2 Mansfield Park to       5475 160460     2      0.0341  
## 3 Mansfield Park and      5438 160460     3      0.0339  
## # ... with 4.038e+04 more rows
```

# Zipf's Law

- Zipf's law is often visualized in log-scale.

```
freq_by_rank %>%  
  ggplot(aes(rank, `term frequency`, color = book)) +  
  geom_line(size = 1.1, alpha = 0.8, show.legend = FALSE) +  
  scale_x_log10() + scale_y_log10()
```



## The bind\_tf\_idf function

- tf-idf finds the words that are important (i.e., common) in a text, but not too common.

```
book_words <- book_words %>%  
  bind_tf_idf(word, book, n)  
print(book_words, n = 5)
```

```
## # A tibble: 40,379 x 7  
##   book      word      n  total    tf    idf  tf_idf  
##   <fct>    <chr> <int> <int>  <dbl> <dbl> <dbl>  
## 1 Mansfield Park the      6206 160460 0.0387      0      0  
## 2 Mansfield Park to       5475 160460 0.0341      0      0  
## 3 Mansfield Park and      5438 160460 0.0339      0      0  
## 4 Emma      to       5239 160996 0.0325      0      0  
## 5 Emma      the      5201 160996 0.0323      0      0  
## # ... with 4.037e+04 more rows
```

- tf-idf are zero for extremely common words.

# The bind\_tf\_idf function

- ▶ Let's look at terms with high tf-idf in Jane Austen's works.

```
book_words %>%  
  select(-total) %>%  
  arrange(desc(tf_idf))
```

```
## # A tibble: 40,379 x 6
```

##	book	word	n	tf	idf	tf_idf
##	<fct>	<chr>	<int>	<dbl>	<dbl>	<dbl>
##	1 Sense & Sensibility	elinor	623	0.00519	1.79	0.00931
##	2 Sense & Sensibility	marianne	492	0.00410	1.79	0.00735
##	3 Mansfield Park	crawford	493	0.00307	1.79	0.00551
##	4 Pride & Prejudice	darcy	373	0.00305	1.79	0.00547
##	5 Persuasion	elliot	254	0.00304	1.79	0.00544
##	6 Emma	emma	786	0.00488	1.10	0.00536
##	7 Northanger Abbey	tilney	196	0.00252	1.79	0.00452
##	8 Emma	weston	389	0.00242	1.79	0.00433
##	9 Pride & Prejudice	bennet	294	0.00241	1.79	0.00431
##	10 Persuasion	wentworth	191	0.00228	1.79	0.00409
##	... with 40,369 more rows					

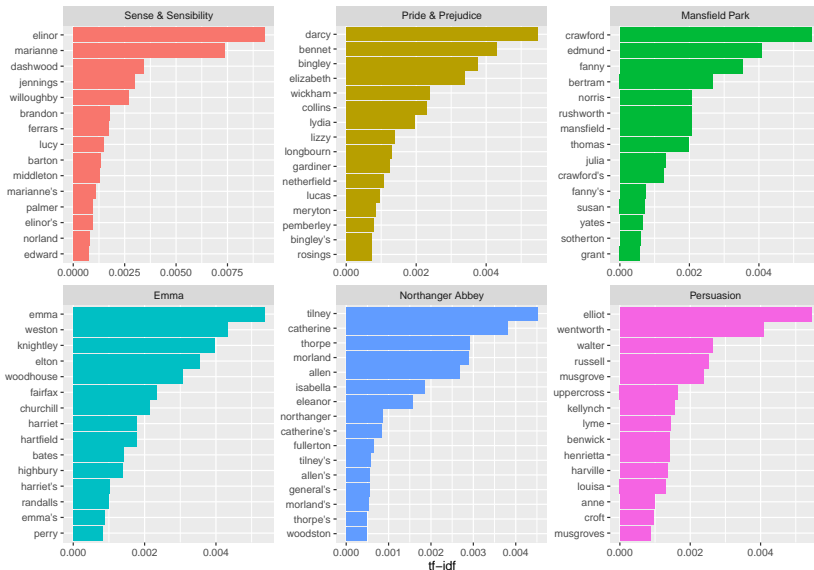


## The bind\_tf\_idf function

```
book_words %>%  
  arrange(desc(tf_idf)) %>%  
  mutate(word = factor(word, levels = rev(unique(word)))) %>%  
  group_by(book) %>%  
  top_n(15) %>%  
  ungroup() %>%  
  ggplot(aes(word, tf_idf, fill = book)) +  
  geom_col(show.legend = FALSE) +  
  labs(x = NULL, y = "tf-idf") +  
  facet_wrap(~book, ncol = 2, scales = "free") +  
  coord_flip()
```

# The bind\_tf\_idf function

## ► Let's visualize!



## A corpus of physics texts

- ▶ Let's download some classic physics texts from Project Gutenberg and see what terms are important in these works, as measured by tf-idf.
  - ▶ Discourse on Floating Bodies by Galileo Galilei
  - ▶ Treatise on Light by Christiaan Huygens
  - ▶ Experiments with Alternate Currents of High Potential and High Frequency by Nikola Tesla
  - ▶ Relativity: The Special and General Theory by Albert Einstein.

```
physics <- gutenbergl_download(c(37729, 14725, 13476, 30155),
                               meta_fields = "author")
```

## Raw counts

```
physics_words <- physics %>%  
  unnest_tokens(word, text) %>%  
  count(author, word, sort = TRUE)  
  
print(physics_words, n = 5)
```

```
## # A tibble: 12,671 x 3  
##   author          word      n  
##   <chr>          <chr> <int>  
## 1 Galilei, Galileo the      3760  
## 2 Tesla, Nikola  the      3604  
## 3 Huygens, Christiaan the      3553  
## 4 Einstein, Albert the      2993  
## 5 Galilei, Galileo of       2049  
## # ... with 1.267e+04 more rows
```

## tf-idf

```
plot_physics <- physics_words %>%  
  bind_tf_idf(word, author, n) %>%  
  mutate(word = fct_reorder(word, tf_idf)) %>%  
  mutate(author = factor(author, levels = c("Galilei, Galileo",  
                                             "Huygens, Christiaan",  
                                             "Tesla, Nikola",  
                                             "Einstein, Albert")))  
  
print(plot_physics, n = 5)
```





```
## # A tibble: 12,671 x 6
```

##	author	word	n	tf	idf	tf_idf
##	<fct>	<fct>	<int>	<dbl>	<dbl>	<dbl>
## 1	Galilei, Galileo	the	3760	0.0935	0	0
## 2	Tesla, Nikola	the	3604	0.0913	0	0
## 3	Huygens, Christiaan	the	3553	0.0928	0	0
## 4	Einstein, Albert	the	2993	0.0952	0	0
## 5	Galilei, Galileo	of	2049	0.0510	0	0
## #	... with 1.267e+04 more rows					

## tf-idf

```
plot_physics %>%  
  group_by(author) %>%  
  top_n(15, tf_idf) %>%  
  ungroup() %>%  
  mutate(word = reorder(word, tf_idf)) %>%  
  ggplot(aes(word, tf_idf, fill = author)) +  
  geom_col(show.legend = FALSE) +  
  labs(x = NULL, y = "tf-idf") +  
  facet_wrap(~author, ncol = 2, scales = "free") +  
  coord_flip()
```

# tf-idf



tf-idf

## tf-idf matrix

```
plot_physics %>% select(author, word, tf_idf) %>%  
  dcast(word ~ author) %>% as_tibble() %>%  
  rename(GG = "Galilei, Galileo", HC = "Huygens, Christiaan"  
    TN = "Tesla, Nikola", EA = "Einstein, Albert")
```

```
## Using tf_idf as value column: use value.var to override.
```

```
## # A tibble: 8,068 x 5
```

```
##   word      GG      HC      TN      EA  
##   <fct> <dbl> <dbl> <dbl> <dbl>  
## 1 1      0      0      0      0  
## 2 10     0      0      0      0  
## 3 11     0      0      0      0  
## 4 12     0      0      0      0  
## 5 13     0      0      0      0  
## 6 14     0      0      0      0  
## 7 15     0      0      0      0  
## 8 16     0      0      0      0  
## 9 17     0      0      0      0  
## 10 18    0      0      0      0  
## # ... with 8,058 more rows
```



# Summary

- ▶ Using term frequency and inverse document frequency allows us to find words that are characteristic for one document in corpus.
- ▶ Exploring term frequency gives us insight into how language is used in a collection of natural language, and `dplyr` verbs like `count()` and `rank()` give us tools to reason about term frequency.